

Practical Part

Note : results, and plots are shown in the colab notebook. 

Task 1

Compare final results on validation set (without the sentences I added) 

| | RNN | GRU |
|----------------|--------|--------|
| Accuracy score | 0.746 | 0.85 |
| F1 score | 0.7687 | 0.8414 |
| ROC AUC score | 0.8488 | 0.9233 |

Compare final results on test set (the sentences I added)

| | RNN | GRU |
|----------------|--------|--------|
| Accuracy score | 0.6667 | 0.8333 |
| F1 score | 0.75 | 0.8571 |
| ROC AUC score | 0.625 | 1.0 |

RNN

1. Hidden-state dimension = 128 (best result for rnn)

Results of Validation set:

Accuracy score: 0.746

F1 score: 0.7687

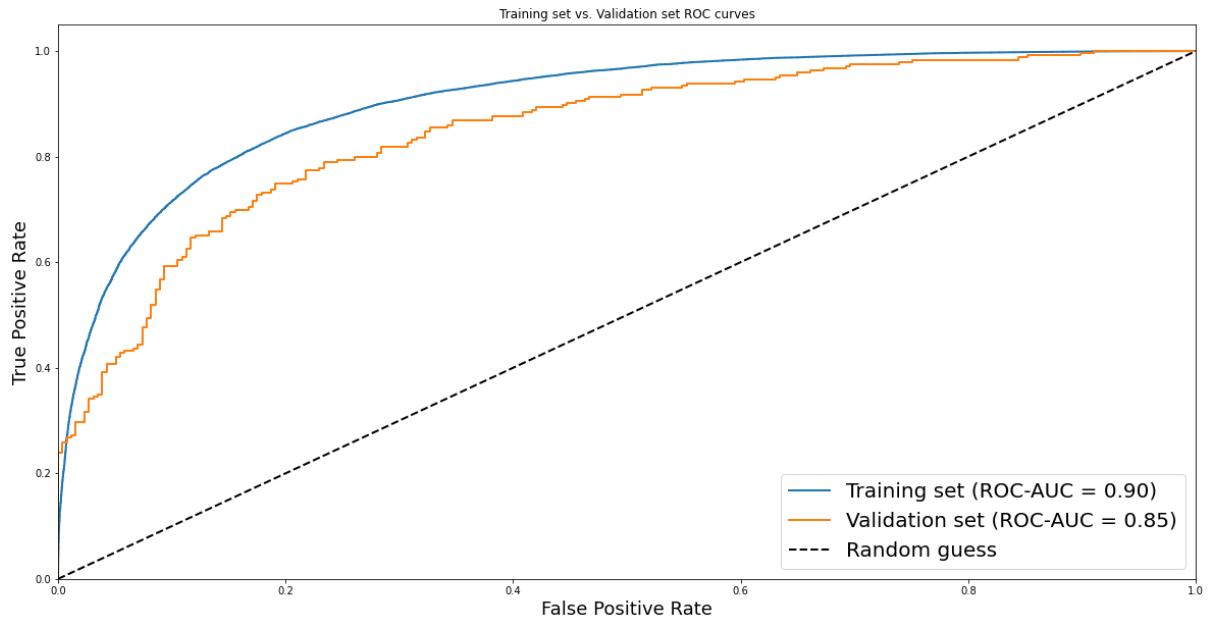
ROC AUC score: 0.8488

Results of Test set:

Accuracy score: 0.6667

F1 score: 0.75

ROC AUC score: 0.625



Example of my test set:

| | Predictions | | True Labels |
|---|-------------|----------|-------------|
| ['This', 'movie', 'is', 'amazing'] | 0.59350574 | Positive | Positive |
| ['I', 'cant', 'recommend', 'this', 'movie', 'highly', 'enough'] | 0.44046244 | Negative | Positive |
| ['I', 'loved', 'this', 'one'] | 0.5743084 | Positive | Positive |
| ['Boring', 'not', 'worth', 'the', 'time'] | 0.25164303 | Negative | Negative |
| ['I', 'had', 'expectations', 'I', 'thought', 'it', 'would', 'be', 'interesting', 'but', 'I', 'was', 'disappointed'] | 0.6438073 | Positive | Negative |

2. Hidden-state dimension = 64

Results of Validation set:

Accuracy score: 0.752

F1 score: 0.7678

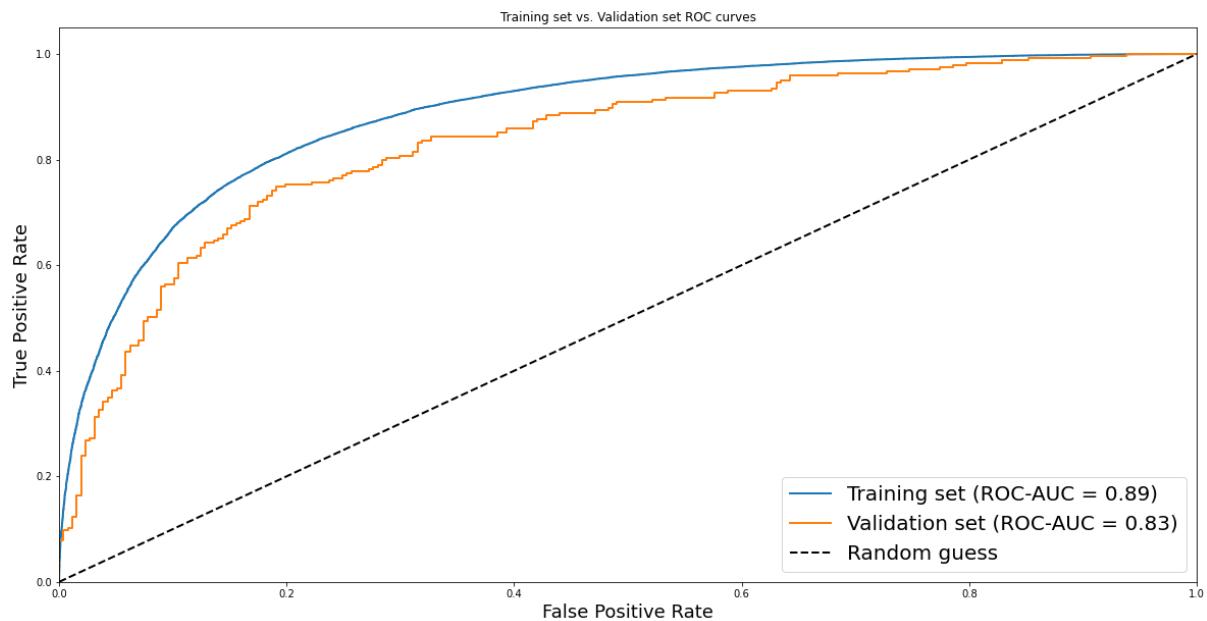
ROC AUC score: 0.8345

Results of Test set:

Accuracy score: 0.5

F1 score: 0.5714

ROC AUC score: 0.625



Example Predictions of the model with highest accuracy:

| | Predictions | | True Labels |
|---|-------------|----------|-------------|
| ['This', 'movie', 'is', 'amazing'] | 0.59350574 | Positive | Positive |
| ['I', 'cant', 'recommend', 'this', 'movie', 'highly', 'enough'] | 0.41895914 | Negative | Positive |
| ['I', 'loved', 'this', 'one'] | 0.4987315 | Negative | Positive |
| ['Boring', 'not', 'worth', 'the', 'time'] | 0.29164532 | Negative | Negative |
| ['I', 'had', 'expectations', 'I', 'thought', 'it', 'would', 'be', 'interesting', 'but', 'I', 'was', 'disappointed'] | 0.5592504 | Positive | Negative |

GRU

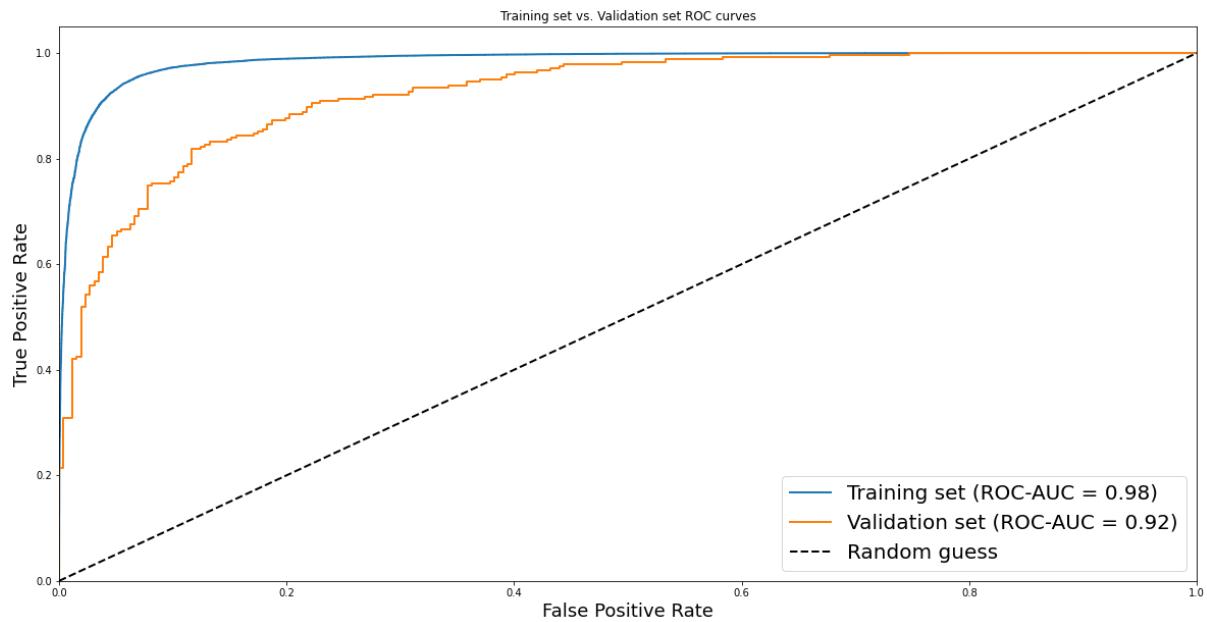
Hidden-state dimension = 64

Results of Validation set:

Accuracy score: 0.85

F1 score: 0.8414

ROC AUC score: 0.9233



Results of Test set:

Accuracy score: 0.8333

F1 score: 0.8571

ROC AUC score: 1.0

Example Predictions:

| | Predictions | | True Labels |
|---|-------------|----------|-------------|
| ['This', 'movie', 'is', 'amazing'] | 0.97527677 | Positive | Positive |
| ['I', 'cant', 'recommend', 'this', 'movie', 'highly', 'enough'] | 0.00666757 | Negative | Positive |
| ['I', 'loved', 'this', 'one'] | 0.9776049 | Positive | Positive |
| ['Boring', 'not', 'worth', 'the', 'time'] | 0.00479398 | Negative | Negative |
| ['I', 'had', 'expectations', 'I', 'thought', 'it', 'would', 'be', 'interesting', 'but', 'I', 'was', 'disappointed'] | 0.00661436 | Negative | Negative |

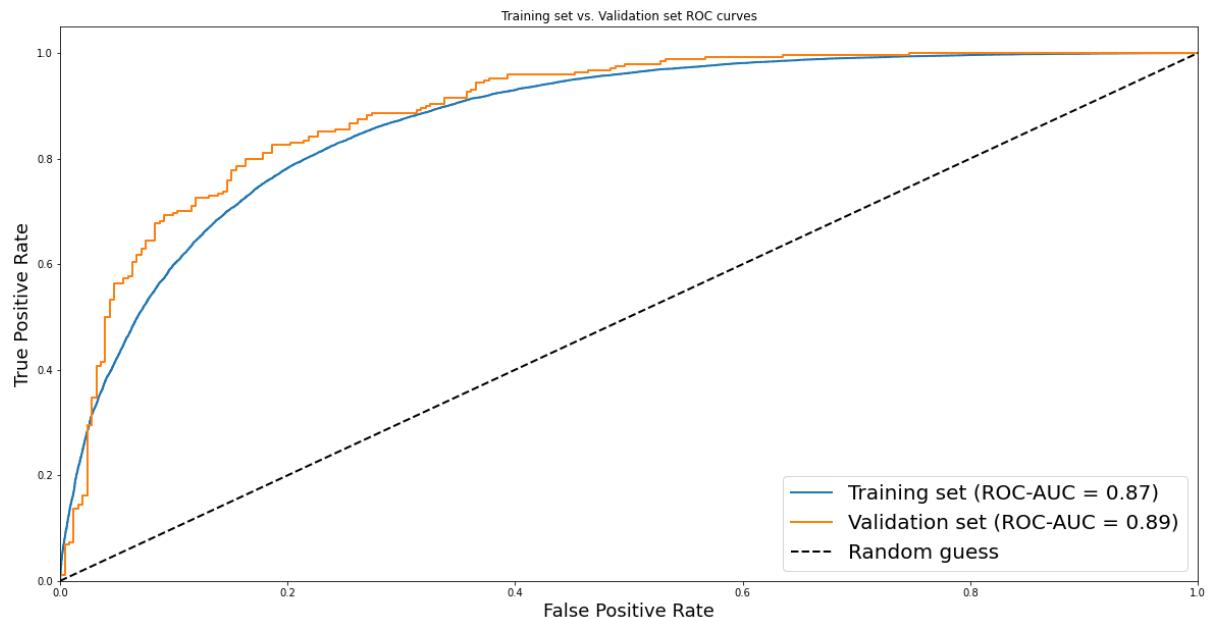
Task 2

Results of Validation set:

Accuracy score: 0.808

F1 score: 0.8132

ROC AUC score: 0.8941



Results of Test set:

Accuracy score: 0.8333

F1 score: 0.8889

ROC AUC score: 0.875

Example Predictions:

| | Predictions | |
|---------|------------------------|----------|
| This | 0.0004223161959089339 | positive |
| movie | 1.209144920721883e-05 | |
| is | 0.001191847026348114 | |
| amazing | 1.6441930711152963e-05 | |

| | Predictions | |
|-----------|------------------------|-----------------|
| I | 1.6565861642447999e-06 | Positive |
| cant | 6.864743045298383e-05 | |
| recommend | 3.7453403933795926e-08 | |
| this | 0.00013868676614947617 | |
| movie | 0.00010089482384501025 | |
| highly | 0.00014803143858443946 | |
| enough | 2.287507605558403e-08 | |

| | Predictions | |
|-------|------------------------|-----------------|
| I | 0.07947801053524017 | Positive |
| loved | 3.2623276638332754e-05 | |
| this | 3.7453403933795926e-08 | |
| one | 0.0016793844988569617 | |

| | Predictions | |
|--------|-----------------------|-----------------|
| boring | 1.209144920721883e-05 | Negative |
| not | 5.591098783952475e-07 | |
| worth | 7.018509495537728e-05 | |
| the | 0.0017538771498948336 | |
| time | 1.94239619304426e-05 | |

| | Predictions | |
|--------------|------------------------|--|
| I | 4.3294221541145816e-05 | |
| had | 0.029746511951088905 | |
| expectations | 0.0005097261164337397 | |
| I | 0.0006210884894244373 | |
| thought | 3.7453403933795926e-08 | |
| it | 0.0005097261164337397 | |
| would | 0.0006115768919698894 | |
| be | 4.483520024223253e-05 | |
| interesting | 3.3501869438623544e-06 | |
| but | 0.00035350158577784896 | |
| I | 0.0005097261164337397 | |
| was | 0.00014803143858443946 | |
| disappointed | 0.0008100144332274795 | |

Positive (wrong prediction)

Consider the sentence : 'I had expectations, I thought it would be interesting, but I was disappointed'

In this example we noticed that this model classified the last sentence as positive review, although it's a negative one. The reason for this mistake could be the misunderstanding of the context. There are words like 'interesting' or 'expectations' that have positive meaning but when reading the words before and after, we can understand that actually the meaning is negative - the movie is not interesting, and the reviewer was disappointed.

Task 3

Let's compare task 2 and task 3 results :

| | Task 2 | Task 3 |
|----------------|--------|--------|
| Accuracy score | 0.8333 | 1.0 |
| F1 score | 0.8889 | 1.0 |
| ROC AUC score | 0.875 | 1.0 |

Now the final layer produces 2 values for each word and uses the additional scalar as a weight for summing the sub scores. This improved the test accuracy since this way the

model can choose which sub score is more important. If a word is important, its weight is high and it will impact the final score of the whole sentence. This score affects the prediction that should be above 0.5 to be considered as positive.

For example the sentence : “This movie is amazing” is a positive review.
The word ‘amazing’ is the most important and its weight is higher compared to task 2.

| | Task 2 | Task 3 |
|----------------|--------------------|-------------------|
| this | 0.00042231 | 0.94605890 |
| movie | 0.00001209 | 1.82377719 |
| is | 0.00119184 | 1.04864975 |
| amazing | 0.000064419 | 0.00756620 |

The sentence “I love this one” is a positive review.
The word ‘loved’ is the most important and its weight is higher compared to task 2.

| | Task 2 | Task 3 |
|--------------|-------------------|--------------------|
| I | 0.07947801 | 1.192771032 |
| loved | 0.00003262 | 0.030170877 |
| this | 3.745340e-08 | 1.992330312 |
| one | 0.0016793 | 0.400458993 |

The sentence “I thought it would be interesting but i was disappointed” is a negative review but task 2 marked this as a positive review.
The word ‘disappointed’ is the most important and its weight is higher compared to task 2.

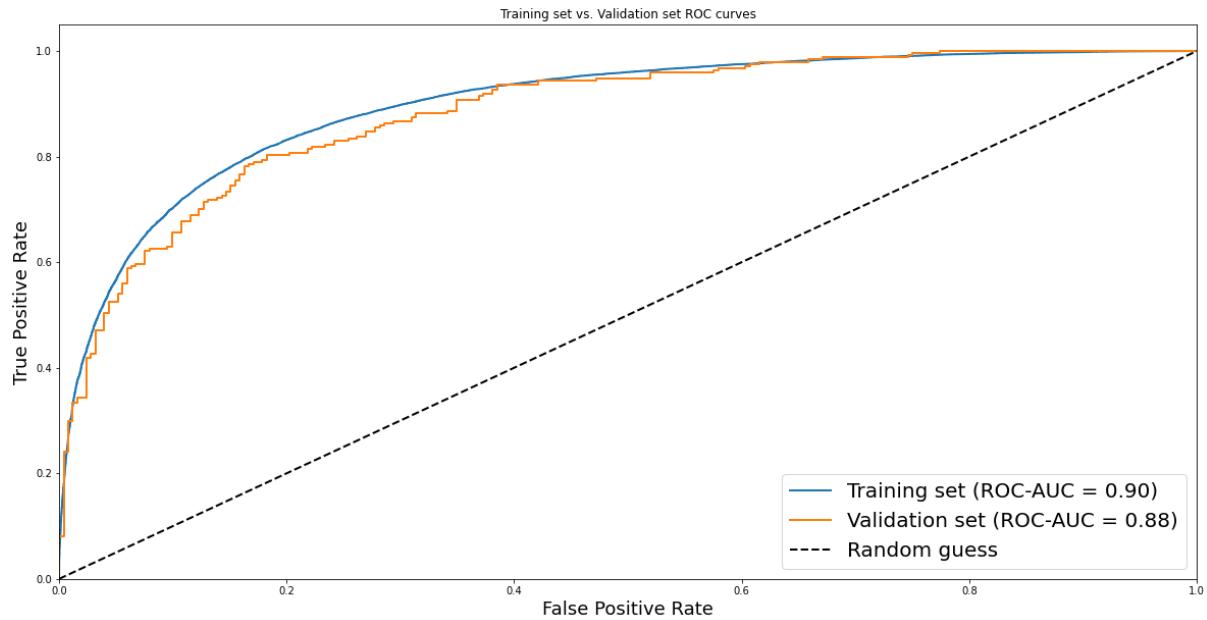
| | Task 2 | Task 3 |
|---------------------|--------------------------|------------------------|
| disappointed | 0.000810014433227 | 1.1151685118675 |

Results of Validation set:

Accuracy score: 0.806

F1 score: 0.8016

ROC AUC score: 0.884



Results of Test set:

Accuracy score: 1.0

F1 score: 1.0

ROC AUC score: 1.0

Example Predictions:

| | Predictions | |
|---------|----------------------|----------|
| This | 0.9460589078080375 | positive |
| movie | 1.823777198791504 | |
| is | 1.0486497581005096 | |
| amazing | 0.007566207583295181 | |

| | Predictions | |
|-----------|--------------------|----------|
| I | 1.2799495458602905 | Positive |
| cant | 1.835557758808136 | |
| recommend | 1.9923303127288818 | |
| this | 1.011188581585884 | |
| movie | 0.8256320059299469 | |
| highly | 0.2529417199548334 | |
| enough | 0.6571245789527893 | |

| | Predictions | |
|-------|----------------------|----------|
| I | 1.1927710324525833 | Positive |
| loved | 0.030170877173077315 | |
| this | 1.9923303127288818 | |
| one | 0.4004589939304424 | |

| | Predictions | |
|--------|---------------------|----------|
| boring | 1.823777198791504 | Negative |
| not | 0.46824405749794096 | |
| worth | 0.2516891956979737 | |
| the | 1.1645564138889313 | |
| time | 0.3470552921498893 | |

| | Predictions | |
|--------------|---------------------|----------|
| I | 0.7417919631116092 | Negative |
| had | 0.6914840147946961 | |
| expectations | 0.35614229179918766 | |
| I | 0.46057990880217403 | |
| thought | 1.9923303127288818 | |
| it | 0.35614229179918766 | |
| would | 0.3617224246263504 | |
| be | 1.1369105577468872 | |
| interesting | 1.0108301350846887 | |
| but | 1.67457914352417 | |
| I | 0.35614229179918766 | |
| was | 0.2529417199548334 | |
| disappointed | 1.1151685118675232 | |

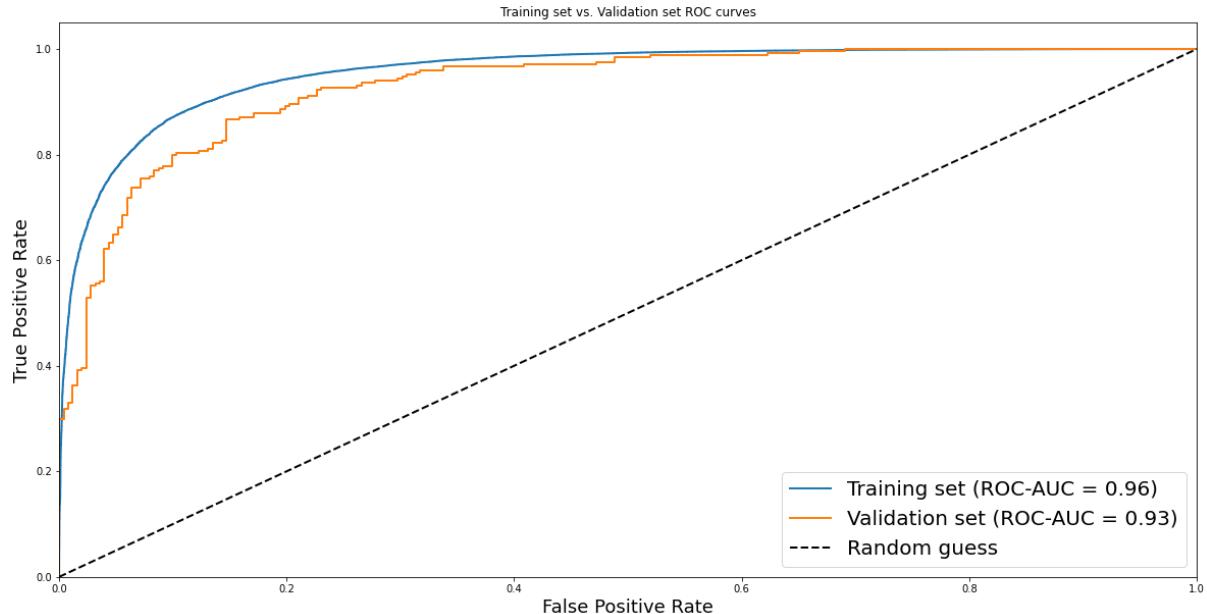
Task 5

Results of Validation set:

Accuracy score: 0.842

F1 score: 0.8384

ROC AUC score: 0.9292



Results of Test set:

Accuracy score: 1.0

F1 score: 1.0

ROC AUC score: 1.0

Example Predictions:

Example Predictions:

| | Predictions | |
|---------|--------------------|----------|
| This | 0.994184136390686 | positive |
| movie | 0.6192785799503326 | |
| is | 0.482318595983088 | |
| amazing | 0.4971533864736557 | |

| | Predictions | |
|-----------|---------------------|-----------------|
| I | 0.31996229849755764 | Positive |
| cant | 0.7433578073978424 | |
| recommend | 0.9995439946651459 | |
| this | 0.6454012244939804 | |
| movie | 1.860028624534607 | |
| highly | 1.079694151878357 | |
| enough | 0.2852265313267708 | |

| | Predictions | |
|-------|--------------------|-----------------|
| I | 0.5464983582496643 | Positive |
| loved | 0.9599781036376953 | |
| this | 0.7846314311027527 | |
| one | 0.2868731766939163 | |

| | Predictions | |
|--------|--------------------|-----------------|
| boring | 1.1153807938098907 | Negative |
| not | 1.3042886853218079 | |
| worth | 1.0846204161643982 | |
| the | 1.4919089674949646 | |
| time | 0.5084910094738007 | |

| | Predictions | |
|--------------|---------------------|--|
| I | 0.19877048209309578 | |
| had | 0.6074221730232239 | |
| expectations | 0.9964616000652313 | |
| I | 0.5945707410573959 | |
| thought | 0.8110296130180359 | |
| it | 0.9962018728256226 | |
| would | 0.9965084791183472 | |
| be | 1.0751320719718933 | |
| interesting | 0.29307248070836067 | |
| but | 0.9883694648742676 | |
| I | 0.9561741650104523 | |
| was | 1.2249450087547302 | |
| disappointed | 1.0710612535476685 | |

Negative

Discussion :

Attention Layer

1. Explain the results - how they differ from before.

Let's compare task 3 and task 5 results :

| | FC sub prediction score | FC sub prediction score + Attention layer |
|----------------|-------------------------|---|
| Accuracy score | 0.806 | 0.842 |
| F1 score | 0.8016 | 0.8384 |
| ROC AUC score | 0.884 | 0.9292 |

As we can see the model accuracy score improved. Moreover when analyzing the prediction score per word , we can see that the Attention layer affects the subscore of close words.

For example the sentence : 'This movie is amazing'
Review is positive and predicted as positive.

| | This | movie | is | amazing |
|--------|-------|-------|-------|---------|
| Task 3 | 0.946 | 1.823 | 1.049 | 0.007 |
| Task 5 | 0.994 | 0.619 | 0.482 | 0.497 |

The self-attention layer we added enables us to learn the correlation between the current words and the previous part of the sentence. Notice that the scores of words that have a relation are close.

For example:

Movie - which movie? this
 Is - who? this movie
 Amazing - who? this movie is

An apparent disadvantage of the model from task 3 is the incapability of the system to remember longer sequences. It could forget the earlier parts of the sequence once it has processed the entire sequence. The attention mechanism resolves this problem.

2. Describe the main principle difference in the predictions abilities that this layer adds to the network and how it can be seen in the results.

The attention model replicates the human attention mechanism and allows the network to focus on a single word from a sentence.

Self-attention is an attention mechanism relating different positions of a single sequence in order to compute a representation of the same sequence. It can help us to learn a correlation between the current words and the previous part of the sentence.

The sub score of each word is computed by measuring the distance between a query to each key. The network extracts interesting features from a sentence, if the distance between words is small , then they get high probability.

Deep Learning Ex 2

$$\frac{\partial}{\partial x} f(x+y, 2x, z) = \frac{\partial}{\partial x} f(u, v, q) = \frac{du}{dx} g(x, y) \cdot D_1 f + \frac{d}{dx} h(x) \cdot D_2 f$$

$$+ \frac{d}{dx} i(z) \cdot D_3 f$$

$$\left. \begin{array}{l} u = g(x, y) = x+y \\ v = h(x) = 2x \\ q = i(z) = z \end{array} \right\}$$

$$D_1 f = \frac{\partial f}{\partial u}, \quad D_2 f = \frac{\partial f}{\partial v}, \quad D_3 f = \frac{\partial f}{\partial q}$$

$$= 1 \cdot D_1 f + 2 \cdot D_2 f$$

(W)

$$f'_1, \dots, n(x) = f'_1 \circ f'_2 \circ f'_3 \circ \dots \circ f'_{n-1} \circ f'_n$$

"g/c i > j p/c g/c

$$f'_1, \dots, j(x) = f'_j$$

$$f'_1, \dots, i(x) = f'_i$$

$$f'_1, \dots, n(x) = f'_1(f_2, \dots, n(x)) \cdot f'_2(f_3, \dots, n(x)) \cdots f'_{n-1}(f_n, n(x))$$

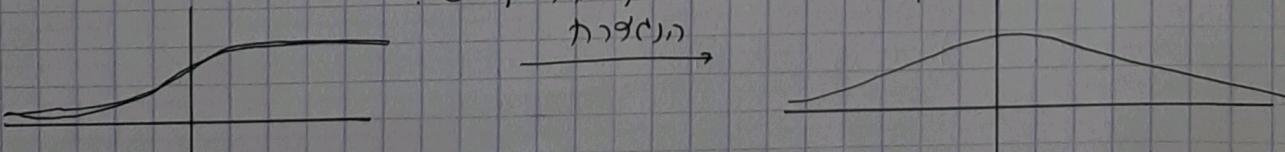
$$\cdot f_n'(x)$$

$$= \prod_{k=1}^n f'_k(f_{k+1}, \dots, n(x))$$

$$f'_{1,2,3}(x) = \frac{f'_1(f_{2,3}(x)) \cdot f'_2(f_3(x))}{f'_3(x)} ; n=3 \text{ 例題 1013}$$

ב) סיבוב מושגים ופונקציית האפסה הינה מוגדרת כפונקציית גראדיאנט גראדיאנט גראדיאנט - Vanishing Gradient גראדיאנט גראדיאנט גראדיאנט דנסט הוא מינימיזציה של פונקציית האפסה ומשתמש באלגוריתם סטנדרטי ומאובטח נאש קומביין עם אלגוריתם סטנדרטי כמו ראנפוקס או ראנפוקס קומביין וראנפוקס גראדיאנט דנסט מושג ערך נמוך יותר מאשר אלגוריתם סטנדרטי.

כטב ררנץ' סיגמוד פונקציית חישובית סigmoid (vanishing gradient)



פונקציית ReLU היא פונקציה לא-ליניארית. בReLU הערך של הפונקציה שווה לאפס אם הערך של ה-
sigmoid הוא מושג כערך נרחב. סימבולי, אם x הוא הערך של ה-
ReLU, אז $f(x) = \max\{0, x\}$. פונקציית ReLU היא פונקציית גיבוב.
למשל, אם x הוא הערך של ה-ReLU, אז $f(x) = \max\{0, x\}$.

$f_1(x, f_2(x, f_3(\dots f_{n-1}(x, f_n(x))))$ (by rule for higher order) (3)

$$\frac{u_1}{v_1} = \frac{x}{f_2(x, f_3(\dots f_{n-1}(x, f_n(x))))} \quad \text{no} \quad *$$

$$\Rightarrow \frac{df_1}{dx} = \frac{du}{dx} \cdot D_1 f_1 + \frac{dv}{dx} f_2(x, f_3(\dots f_{n-1}(x, f_n(x)))) \cdot D_2 f_1$$

$$D_1 f_1 = \frac{\partial f_1}{\partial u_1}, \quad D_2 f_1 = \frac{\partial f_1}{\partial v_1} \quad \text{no} \quad *$$

$$u_2 = x, \quad v_2 = f_3(\dots f_{n-1}(x, f_n(x))) \quad \text{no} \quad *$$

$$D_1 f_2 = \frac{\partial f_2}{\partial u_2}, \quad D_2 f_2 = \frac{\partial f_2}{\partial v_2} \quad \text{no} \quad *$$

$$\Rightarrow \frac{df_2}{dx} = \frac{du}{dx} \cdot D_1 f_2 + \frac{dv}{dx} f_3(\dots f_{n-1}(x, f_n(x)))$$

$$\frac{d}{dx} f_1(x, f_2(x, f_3(\dots f_{n-1}(x, f_n(x))))) \quad \text{no} \quad \text{no}$$

$$= 1 \cdot \underline{D_1 f_1} + \underline{D_2 f_1} \cdot \left[1 \cdot \underline{D_1 f_2} + f_3(\dots f_{n-1}(x, f_n(x))) \cdot \dots \cdot f_{n-1}'(x, f_n(x)) \cdot (1 \cdot \underline{D_1 f_{n-1}} + f_n'(x) \cdot \underline{D_2 f_{n-1}}) \right] \approx 0$$

Skip Connection : $f(x + g(h(x)))$

(5)

Vanishing - \rightarrow skip connection מושך אפס. Skip connection נזק אם ה- g מושך אפס. skip connection מושך אפס אם $g'(x+h(x)) = 0$. skip connection מושך אפס אם $g'(x+h(x)) \cdot h'(x) = 0$.

$$\frac{\partial}{\partial x} f(g(h(x))) = f'(g(h(x))) \cdot g'(h(x)) \cdot h'(x)$$

כדי שskip connection מושך אפס, צריך ש- $h'(x) = 0$.

בנוסף $f(x + g(x + h(x)))$ מושך אפס.

$$\begin{aligned} \frac{\partial}{\partial x} f(x + g(x + h(x))) &= f'(x + g(x + h(x))) \cdot (1 + g'(x + h(x))) \\ &\quad \cdot (1 + h'(x)) \\ &= (f'(x + g(x + h(x))) + f'(x + g(x + h(x))) \cdot \\ &\quad g'(x + h(x))) (1 + h'(x)) \\ &= f'(x + g(x + h(x))) + f'(x + g(x + h(x))) \cdot g'(x + h(x)) \\ &\quad + f'(x + g(x + h(x))) \cdot h'(x) \\ &\quad + f'(x + g(x + h(x))) \cdot g'(x + h(x)) \cdot h'(x) \end{aligned}$$

skip connection מושך אפס אם $h'(x) = 0$. skip connection מושך אפס אם $g'(x + h(x)) = 0$. skip connection מושך אפס אם $g'(x + h(x)) \cdot h'(x) = 0$. skip connection מושך אפס אם $g'(x + h(x)) \cdot g'(x + h(x)) = 0$.

ב- $f(x + g(x + h(x)))$, back propagation מושך אפס אם $g'(x + h(x)) = 0$.

אנו מודים בפונקציית loss על skip connection.

Speech Recognition - Audio to text (IC)

CNN מושך אפס אם x מושך אפס. CNN מושך אפס אם x מושך אפס.

CNN מושך אפס אם x מושך אפס. CNN מושך אפס אם x מושך אפס. CNN מושך אפס אם x מושך אפס. CNN מושך אפס אם x מושך אפס. CNN מושך אפס אם x מושך אפס.

pooling מושך אפס אם x מושך אפס. pooling מושך אפס אם x מושך אפס. pooling מושך אפס אם x מושך אפס.

Answer Question

(2)

many to many RNN have many hidden states, so it's hard to find a good way to do this. One solution is to use attention mechanism.

In this, we can use GRU hidden states to calculate which word is more likely to be the next word. This is called Attention layer. It takes the hidden state of the previous word and the hidden state of the current word and produces a score between 0 and 1. This score is then multiplied with the hidden state of the previous word to produce a new hidden state. This new hidden state is then passed through an FC layer to produce the final output.

Sentiment Analysis

(3)

In this, we can use GRU hidden states to calculate the sentiment of the sentence. We can use an FC layer to produce a score between 0 and 1. This score is then multiplied with the hidden state of the previous word to produce a new hidden state. This new hidden state is then passed through an FC layer to produce the final output.

Q. How does attention work? Answer questions by writing code in Python.

A. In this, we can use GRU hidden states to calculate the sentiment of the sentence. We can use an FC layer to produce a score between 0 and 1. This score is then multiplied with the hidden state of the previous word to produce a new hidden state. This new hidden state is then passed through an FC layer to produce the final output.

Q. What is vanishing gradient? Answer questions by writing code in Python.

A. Vanishing gradient is a problem in RNN where the gradients of the loss function with respect to the hidden states become very small as they pass through the network. This is because the gradients are multiplied by the weights of the hidden states at each time step. If the weights are small, the gradients will also be small. This can lead to slow learning and even stop learning completely. To solve this, we can use a technique called gradient clipping, which limits the size of the gradients. Another solution is to use a different type of RNN, such as LSTM or GRU, which are less prone to vanishing gradients.

Image Classification

(4)

In this, we can use CNN to extract features from the image. These features are then passed through an FC layer to produce the final output. This is called a CNN - FC architecture.

Q. What is translation invariant? Answer questions by writing code in Python.

A. Translation invariant means that the output of the network should be the same regardless of the position of the input. This is achieved by using a convolutional layer with stride 1. The input is first passed through an encoder, which consists of several layers of convolutional and fully connected layers. The output of the encoder is then passed through a decoder, which consists of several layers of convolutional and fully connected layers. The final output is the translation invariant representation of the input.

3. Text - to - image:

a) Describe the architecture of a network that reads a sentence and generates an image based on this text. Do not address the question of how such a network should be trained, just explain why it should have the capacity to perform this task. You can assume that the images come from a restricted class of images e.g. faces and can be encoded (and decoded) in a low-dimensional latent space.

image - n. text-to-image תרגום טקסט לתמונה, צילום מסך, צילום מסך

gradient vanishing problem

layer

ההיבר אוניברסיטאות מוסמך בתקופה ה-23' הנקראת **Hidden Vector**.

פָּרְטָה feature (o feature vector) נויה בז'קצ'ינטס בעקבות מושג אחד או יותר.

עומק תרשים הינה א-סימטרית (ב-Attention layer)

ב-ט'כיד אונדניריג גוט-ווען דלאוון סטראָפּ. אַנְדָּר דִּנְצָה (ט'כיד) - (ט'לט'ו).

הטביה נולאה עז (text to image)

text → \vec{x} column \vec{y} (embedding) image → \vec{y} (captioning)

לפניהם נתקה רג'יסטר וטולר וויליאם סמואל ג'ונס, כוותר יהודים נאזרחים עיר

הנתק (פנורמי) כריזטוטן אם כוון הזרם שוואץ'ם ולבסוף חירך ה-3 כיריך

המודול ימוך את הערך ב- NN_1 . ה-encoder ימוך את הערך ב- NN_2 .

הנ'ויל רודט נאנויאנד ("בז'ן") הרכיך מהאנדרון הפליגרנטאלן ציז'ון.

۱۰۷

הוילט, גנטה ה' ציון ו-200 מטרים פלטאליק. בתקופה זו, בין 8 ו-10 ג'יג'ר נרעו הימנאים כ-N69. היה לנו גם ראייה נוספת משלוחת הדרוזית מ-1999 שראתה קבוצת קבינה יפה נספחה לארכיאולוג שארם סוליבן מילר. מילר מזכיר במאמריו שראתה קבוצה דתנית מושג עתיקות, טהרה את רוג'ה מהר הבית ופיזה נערם מהקדש - רוג'ה אלוקת ירושלים ורומי. עם זאת, מילר מזכיר שראתה קבוצה דתנית מושג עתיקות, טהרה את רוג'ה מהר הבית ופיזה נערם מהקדש - רוג'ה אלוקת ירושלים ורומי. עם זאת, מילר מזכיר שראתה קבוצה דתנית מושג עתיקות, טהרה את רוג'ה מהר הבית ופיזה נערם מהקדש - רוג'ה אלוקת ירושלים ורומי.

ג'י'ע'ן ז' פֿאַלְמָה (עֲמָנָה)

b. Assume that the latent codes are made of 4 vectors that correspond to four quadrants of an image. Explain how an attention layer can be used to allow such a network to better support regional descriptions in the input text. What would be the queries, keys and values?

לכיכר או לתרנגולת כוונת אפקת ה-attention (link) יתבצע במלבוב (linking). על כן, נזכיר

(גנרי, א. כב. 1"ב) ו/או ימי קבניר, צייר (בפ' גנרי א' (גנרי ד' (גנרי א'))

בְּיַעֲנֵה לְתַחַת לִשְׁבָּר אֶלְעָמָן

(d) Form NLP features from word vectors

אנו יכולים לחלק נתונים על ידי feature vectors בלבד. כלומר, אם נזקיף על feature vectors

לעומת מודל ה-Attention, מודל ה-GRUCell מפעיל אטטנציה על כל תimestep בפער.

כטבנ'ו Attention layers

השאלה הינה: מה הערך של hidden vector query אם הhidden layer נוציא

הנתקל בפונקציית `key` ב-`map` ו-`reduce` ו-`groupByKey` ו-`join` מ-`Spark`.

10. גמונת פילס נבנית על ידי מנגנון Softmax נ-ענו!

בכך (וכותם מלהרשותן וה-value יהיה לנו את המונטג'ה (לעומת גיא) ונמצא בפה (בפה) מושג מושג (כגון קד) או אולי אף רמה גאותה (ולא יותר).